

# CFRM 421 Group 8 Final Project

## SPY High-Frequency Trading Strategy Prediction with Machine Learning

Ziyu Liao    Tong Wu    Chenxu Wang  
Max Chen    Yiheng Tong    Yipeng Hu

June 2025

## Abstract

We apply machine learning to predict intraday trading actions—buy, sell, or hold—for SPY ETF using one year of minute-level data. After engineering RSI-based features and reducing dimensionality via PCA, we train and evaluate multiple classifiers, including k-NN, GBDT, and Logistic Regression. A soft-voting ensemble of top models achieves over best accuracy, F1-score, and confusion matrix results. However, the results shows the limitation of our models, highlighting key challenges in high-frequency trading prediction.

# Project Goal and Problem Statement

- ▶ **Goal:** Apply machine learning to predict a trading strategy for SPY. Specifically, forecast the next one-minute action: buy, sell, or hold.
- ▶ **Problem:** Intraday financial data is noisy and high-frequency, making accurate prediction challenging. We frame this as a 3-class classification problem: Buy (+1), Sell (-1), or Hold (0).
- ▶ We aim to design a model that can capture informative patterns from tick-level data to support trading decisions. A successful model could enable a profitable high-frequency trading strategy.
- ▶ Key challenges include feature engineering for high-frequency data, handling class imbalance (as most minute-to-minute changes are small), and avoiding overfitting to noise.

# Data Overview

- ▶ **Dataset:** 1-year of SPY (S&P 500 ETF) tick-level trade data aggregated to 1-minute bars.
- ▶ **Period:** July 1, 2020 – July 1, 2021. Data downloaded from [Finhub](#).
- ▶ **Features per bar:** timestamp, open, high, low, close, volume, dollar\_volume, tick\_count, trade\_size\_mean, trade\_size\_std, zero\_return\_count, price\_direction\_ratio, large\_trade\_count, large\_trade\_volume, vwap, large\_trade\_ratio, large\_trade\_volume\_ratio.
- ▶ **Target variable:** defined based on next-minute trading strategy. See next slide.

# Data Processing & Feature Engineering

- ▶ **Data cleaning:** Filled sporadic missing values via linear interpolation.
- ▶ **Target labeling:** Defined action labels based on forward log return over 5 minutes with a 1% threshold. (This introduced class imbalance, as "Hold" was the majority case.)
- ▶ **Technical indicators:** Computed 14-period Relative Strength Index (RSI) for each unbounded price/volume feature (open, high, low, close, volume, etc.). RSI transforms raw values into a bounded 0–100 range, providing normalized momentum signals.
- ▶ After adding RSI features, we dropped the original unbounded feature columns. This yielded a feature set consisting entirely of bounded indicators (RSIs of price, volume, and trade metrics), along with existing ratio features (e.g., `price_direction_ratio`).

# Data Processing & Feature Engineering

- ▶ **Normalization & PCA:** Standardized features (zero-mean, unit-std). Applied PCA and retained the main components  $n = 7$  (covering the variance  $\geq 90\%$ ) for use in modeling, to reduce dimensionality and noise.
- ▶ **Train/Validation/Test split:** We employed a chronological time-series split method to structure our data into training, validation, and test sets. Each data sample consisted of sequences of 16-minute feature windows, with the subsequent minute's action serving as the target variable. This time-aware splitting ensures a realistic and robust evaluation, effectively preventing look-ahead bias by strictly preserving the temporal order in model training and validation.

## Exploratory Analysis: Clustering Regimes

- ▶ Before modeling, we explored the data structure using unsupervised learning:
  - ▶ Performed PCA on the feature set (after preprocessing) to reduce dimensionality. The top 7 principal components explained over 90% of the variance.
- ▶ Applied K-Means clustering (with  $k = 2$ ) and DBSCAN on the PCA-transformed data to identify potential market regimes or patterns in the one-minute bars.
- ▶ **Result:** The data points formed distinct clusters in principal component space (see next slide), suggesting that the market may exhibit different intraday states (e.g., high-volatility vs. low-volatility periods).

# Cluster Visualization in PCA Space

Clustering visualization (PC 1 \* PC 2)




Figure: Clustering of data in PCA space

## Model Selection and Implementation (Traditional ML)

- ▶ We implemented and systematically tuned a diverse set of supervised learning algorithms to tackle the classification task. Hyperparameters were optimized using GridSearchCV with 5-fold stratified cross-validation:
  - ▶ **Logistic Regression:** Multinomial logistic classifier (softmax) with balanced class weights, tuned  $C = 1.0$ , solver lbgfs,  $L_2$  regularization.
  - ▶ **Ridge Classifier:** Linear classifier with optimized  $L_2$  regularization strength  $\alpha = 0.1$ .
  - ▶ **k-Nearest Neighbors:** Selected  $k = 5$ , Euclidean distance ( $p=2$ ), distance-based weighting scheme for prediction.
  - ▶ **Random Forest:** Optimized ensemble of 200 decision trees, max depth=4, minimal leaf size=1, using  $\sqrt{d}$  features.
  - ▶ **Histogram Gradient Boosted Trees:** Gradient boosting with trees of depth=4, slow learning rate (0.01), 600 boosting iterations, and balanced class weights.
  - ▶ **Multi-Layer Perceptron:** Two-layer feed-forward neural network (128, 64 neurons), ReLU activations, optimized regularization ( $\alpha = 10^{-5}$ ), initial learning rate ( $10^{-4}$ ), early stopping enabled.

## Model Selection and Implementation (Traditional ML)

- ▶ The hyperparameter tuning process utilized a smaller representative subset (first 10,000 time-series samples) of the training data to efficiently identify optimal parameters.
- ▶ Finalized models were then retrained using the full training dataset, and performance was rigorously evaluated on a hold-out test set.
- ▶ **Class imbalance management:** Implemented class-weight balancing explicitly for logistic regression, random forest, and gradient boosting. Models were evaluated primarily using macro-F1 score to ensure balanced performance across classes.
- ▶ **Reproducibility and framework:** Python scikit-learn implementation with fixed random states (`random_state=42`) ensured consistent and reproducible results across runs.
- ▶ **Cross-validation strategy:** Stratified K-Fold (5 splits, randomized) to maintain class distributions during hyperparameter tuning.

## Model Selection and Implementation (Traditional ML)

- ▶ **Rationale for Model Selection:** We selected these diverse models to explore different types of algorithms, each offering distinct predictive capabilities suited to our classification task:
  - ▶ **Logistic Regression and Ridge Classifier** were chosen as robust baseline linear classifiers, offering interpretability and computational efficiency.
  - ▶ **k-Nearest Neighbors (k-NN)** was included for its simplicity and strong capability in capturing local data patterns, especially beneficial in imbalanced settings.
  - ▶ **Random Forest and Histogram Gradient Boosted Trees (GBDT)** provided powerful ensemble-based approaches that excel at handling nonlinear relationships and complex interactions between features.
  - ▶ **Multi-Layer Perceptron (MLP)** offered flexibility in modeling complex, nonlinear relationships and potential interactions in high-dimensional spaces.

## Results: Classification Performance


- ▶ **Evaluation Metrics:** We evaluated models on the test set using overall accuracy and macro-averaged F1 score (macro-F1 treats each class equally, highlighting performance on minority classes).
- ▶ **Accuracy vs. F1:** Many models achieved very high accuracy (near 99% or above) by primarily predicting the majority class (Hold). However, their macro-F1 scores were low ( $\approx 0.33$ ), indicating poor prediction of the minority Buy/Sell classes (similar to always guessing “Hold” ).
- ▶ The k-NN model was a notable exception: it attained macro-F1  $\approx 0.72$ , substantially higher than others. k-NN managed to capture a subset of the minority class instances, giving more balanced performance (while still maintaining  $\sim 100\%$  accuracy).

# Model Performance: Accuracy vs. Macro-F1

We provide detailed metrics to illustrate models' predictive capability. High accuracy alone is misleading due to class imbalance; thus macro-F1 reveals true performance on minority classes.

Model	Acc	F1
LogReg	0.890	0.316
Ridge	1.000	0.333
k-NN	1.000	0.717
RF	1.000	0.333
Hist GBDT	0.922	0.353
MLP	1.000	0.333

**Table:** Test set performance.  
Macro-F1 reveals minority  
class prediction power.



**Figure:** Visual comparison of model metrics.

## Results: Confusion Matrix Insights

Confusion matrices reveal model-specific prediction patterns clearly:

- ▶ **Logistic Regression:** Identifies most "Buy" signals correctly (75%), but struggles slightly more with "Sell" (25%). Reasonably accurate Hold classification (89%).
- ▶ **k-NN:** Strong at minority class detection—accurately identifies 75% of "Buy", moderate (25%) accuracy on "Sell", and perfect "Hold" prediction (100%).
- ▶ **Ridge Random Forest:** Exclusively predict "Hold", ignoring minority classes entirely (0% minority class accuracy).
- ▶ **Hist GBDT:** Similar to logistic regression—predicts 75% of "Buy", but misclassifies many "Sell" signals (only 25% correctly identified). Holds generally well predicted (92%).
- ▶ **MLP:** Identical issue as Ridge/RF—predicts solely "Hold".

**Next slide:** Visual comparison via normalized confusion matrices.

# Normalized Confusion Matrices - Validation Set

Normalized confusion matrices demonstrate clear distinctions in minority-class prediction accuracy among models.




Figure: Normalized confusion matrices

## Results: Soft Voting Ensemble

- ▶ **Ensemble Method:** Combined the predictions of three well-performing base models:
  - ▶ Logistic Regression (Softmax)
  - ▶ k-Nearest Neighbors (k-NN)
  - ▶ Histogram Gradient Boosting Trees (Hist GBDT)
- ▶ Utilized a `VotingClassifier` with *soft voting*, leveraging predicted class probabilities to form a robust ensemble prediction.
- ▶ **Performance Metrics:**
  - ▶ *Validation Set:*
    - ▶ Accuracy: **99.61%**
    - ▶ Weighted F1-score: **99.78%**
  - ▶ *Test Set (Final Evaluation):*
    - ▶ Accuracy: **99.70%**
    - ▶ Weighted F1-score: **99.84%**
- ▶ **Key Insights:**
  - ▶ Ensemble significantly outperforms individual models in minority class detection and overall predictive stability.
  - ▶ Near-perfect accuracy and very high weighted F1-score demonstrate superior generalization capability and balanced performance across all classes.

# Confusion matrix of Soft Voting Classifier




Figure: Normalized results of Voting classifier on valid and test set

# Results: Backtesting Trading Strategy (Visualization)




Figure: Cumulative PnL (out-of-sample) for each model vs. Buy & Hold.

# Conclusion and Key Findings

- ▶ Built an end-to-end ML framework for high-frequency action prediction using feature engineering (e.g. RSI), diverse algorithms, unsupervised learning.
- ▶ **Key Findings:**
  - ▶ Extreme class imbalance was a major challenge — most models predicted the majority class (no-action), achieving high accuracy but missing rare buy/sell signals.
  - ▶ The k-NN model performed relatively better in detecting such signals, but predictive power for minority classes remained limited.
  - ▶ Out-of-sample simulations showed that even our best model failed to yield positive risk-adjusted returns.
  - ▶ **Insight:** High accuracy alone doesn't guarantee profit, especially when driven by trivial predictions. Metrics like F1 are crucial for evaluation.

# Lessons and Future Work

- ▶ **Lessons:**
  - ▶ Short-term market prediction is inherently difficult but doable.
  - ▶ Accuracy must be aligned with meaningful signal detection.
- ▶ **Future Improvements:**
  - ▶ Refine labeling scheme to address class imbalance.
  - ▶ Incorporate regime detection to switch strategies under different market conditions.
  - ▶ Explore online learning and RL to adapt over time.
  - ▶ Add richer timestamps and features (e.g. order book data, news sentiment) to improve signal-to-noise ratio.
- ▶ The project offered hands-on experience and revealed both potential and challenges in applying ML to finance.

# Appendix

- ▶ **Video Presentation:** Link to Google Drive
- ▶ **Jupyter Notebook:** Link to Github
- ▶ **Data Source:** Link to Google Drive